SAME-BENCHMARK CANDIDATE GALLERIES · 40 TASKS

## Human query → Agent gallery

A

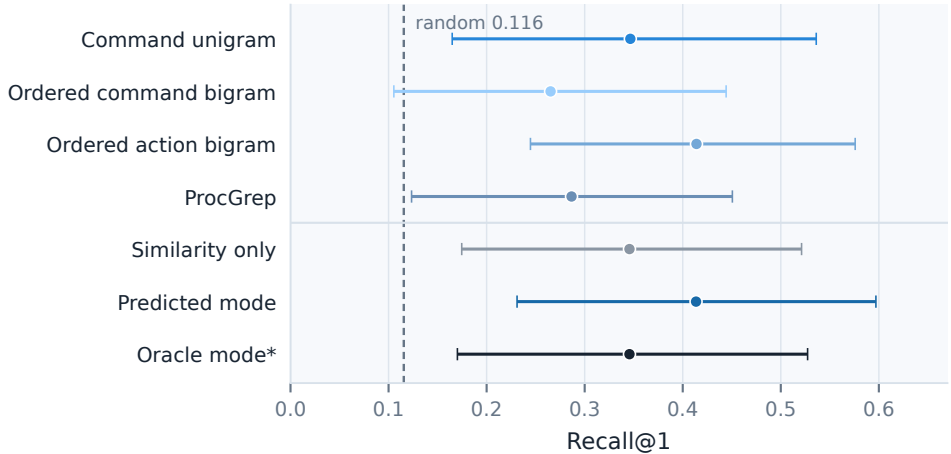


| | Estimate [95% task CI] |
|---|---|
| Command unigram | 0.347 [0.165, 0.536] |
| Ordered command bigram | 0.265 [0.105, 0.444] |
| Ordered action bigram | 0.414 [0.245, 0.576] |
| ProcGrep | 0.287 [0.123, 0.451] |
| Similarity only | 0.346 [0.175, 0.521] |
| Predicted mode | 0.414 [0.231, 0.597] |
| Oracle mode* | 0.346 [0.170, 0.527] |

Whiskers show 95% task-cluster intervals. Oracle mode uses ground-truth labels at inference.